## Stat 432 HW 02

Name: Ahmadreza Eslaminia, netID: ae15

#### Summer 2024

Include the R code for this HW.

```
knitr::opts_chunk$set(echo = TRUE)
library(ISLR2)
library(GGally)
```

There are some useful R chunk options that you may use (for this entire semester):

- echo Display code in output document (default = TRUE)
- include Include chunk in document after running (default = TRUE)
- $\bullet$  message display code messages in document (default = TRUE)
- results (default = 'markup')
  - 'asis' passthrough results
  - 'hide' do not display results
  - 'hold' put all results below all code
- $\bullet$  error Display error messages in doc (TRUE) or stop render when errors occur (FALSE) (default = FALSE)

See R markdown cheat sheet for more information.

#### Question 1 (Linear Regression)

We have N observations of  $(X_1, X_2, \dots, X_p, Y)$ .

Let us use the following notations:

- X is a the  $N \times (p+1)$  matrix with each row as an input vector (with a 1 in the first position),
- $\mathbf{y}$  be the N-vector of outputs and

$$\bullet \quad \boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix}.$$

Then we may write the multiple linear regression model as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \epsilon$$
.

Show that

$$\widehat{\boldsymbol{\beta}} = (\mathbf{X}^{\mathrm{T}}\mathbf{X})^{-1}\mathbf{X}^{\mathrm{T}}\mathbf{y}.$$

minimizes RSS.

Answer: for the linear regression model:

$$\mathbf{v} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

where: -  $\mathbf{X}$  is an  $N \times (p+1)$  matrix of inputs. -  $\mathbf{y}$  is an N-vector of outputs. -  $\boldsymbol{\beta}$  is a (p+1)-vector of coefficients.

RSS is defined as:

$$RSS = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

To minimize the RSS set deravative to beta to zero:

$$\frac{\partial RSS}{\partial \boldsymbol{\beta}} = -2\mathbf{X}^T\mathbf{y} + 2\mathbf{X}^T\mathbf{X}\boldsymbol{\beta} = 0$$

$$\mathbf{X}^T \mathbf{X} \boldsymbol{\beta} = \mathbf{X}^T \mathbf{y}$$

Assuming  $\mathbf{X}^T\mathbf{X}$  is invertible, we find the following as the minimizer:

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{v}$$

#### Question 2 (Linear Regression)

This question relates to the College data set, which can be found in the file College.csv. It contains a number of variables for 777 different universities and colleges in the US.

(from the previous HW) Use the read.csv() function to read the data into R. Call the loaded data college. Make sure that you have the directory set to the correct location for the data.

Before moving on, we're not going to use the college name, so you may remove X variable from data.

Also, make sure categorical variables are set as factor variables.

Split your data into two parts: a testing data that contains 100 observations, and the rest as training data. You may use sample function to get the indices of the testing data. For this question, you need to set a random seed while generating this split so that the result can be replicated. Use 4322 as the random seed. Report the mean of Outstate of your testing data and training data, respectively.

```
college <- read.csv("College.csv")
college <- college[ , !(names(college) %in% "X")]

# factorizing the categirical features ( only private)
college$Private <- as.factor(college$Private)
set.seed(4322)
summary(college)</pre>
```

```
##
    Private
                    Apps
                                    Accept
                                                      Enroll
                                                                    Top10perc
                           81
##
    No :212
               Min.
                                Min.
                                        :
                                            72
                                                 Min.
                                                         :
                                                            35
                                                                  Min.
                                                                          : 1.00
    Yes:565
##
               1st Qu.:
                          776
                                1st Qu.:
                                           604
                                                  1st Qu.: 242
                                                                  1st Qu.:15.00
##
               Median: 1558
                                Median: 1110
                                                  Median: 434
                                                                  Median :23.00
##
               Mean
                       : 3002
                                Mean
                                        : 2019
                                                  Mean
                                                         : 780
                                                                  Mean
                                                                          :27.56
##
               3rd Qu.: 3624
                                3rd Qu.: 2424
                                                  3rd Qu.: 902
                                                                  3rd Qu.:35.00
##
               Max.
                      :48094
                                Max.
                                        :26330
                                                  Max.
                                                          :6392
                                                                  Max.
                                                                          :96.00
##
      Top25perc
                      F. Undergrad
                                        P. Undergrad
                                                             Outstate
               9.0
                                                                  : 2340
##
                     Min.
                                139
                                      Min.
                                              :
                                                    1.0
                                                          Min.
##
    1st Qu.: 41.0
                     1st Qu.:
                                992
                                       1st Qu.:
                                                   95.0
                                                          1st Qu.: 7320
    Median: 54.0
                     Median: 1707
                                                  353.0
                                                          Median: 9990
##
                                      Median:
            : 55.8
                             : 3700
##
    Mean
                     Mean
                                       Mean
                                                 855.3
                                                          Mean
                                                                  :10441
##
    3rd Qu.: 69.0
                     3rd Qu.: 4005
                                       3rd Qu.:
                                                 967.0
                                                          3rd Qu.:12925
                             :31643
##
    Max.
            :100.0
                     Max.
                                              :21836.0
                                                          Max.
                                                                  :21700
                                      Max.
##
      Room.Board
                        Books
                                          Personal
                                                            PhD
            :1780
                               96.0
                                              : 250
                                                                  8.00
##
    Min.
                    Min.
                            :
                                      Min.
                                                       Min.
                                                               :
                                                       1st Qu.: 62.00
##
    1st Qu.:3597
                    1st Qu.: 470.0
                                       1st Qu.: 850
##
    Median:4200
                    Median : 500.0
                                      Median:1200
                                                       Median: 75.00
##
    Mean
            :4358
                    Mean
                            : 549.4
                                       Mean
                                              :1341
                                                       Mean
                                                               : 72.66
                    3rd Qu.: 600.0
                                                       3rd Qu.: 85.00
##
    3rd Qu.:5050
                                       3rd Qu.:1700
##
    Max.
            :8124
                            :2340.0
                                              :6800
                                                               :103.00
                    Max.
                                       Max.
                                                       Max.
##
       Terminal
                       S.F.Ratio
                                       perc.alumni
                                                            Expend
                             : 2.50
##
    Min.
            : 24.0
                     Min.
                                       Min.
                                              : 0.00
                                                        Min.
                                                                : 3186
##
    1st Qu.: 71.0
                     1st Qu.:11.50
                                       1st Qu.:13.00
                                                        1st Qu.: 6751
##
    Median: 82.0
                     Median :13.60
                                      Median :21.00
                                                        Median: 8377
##
            : 79.7
                             :14.09
                                              :22.74
                                                                : 9660
    Mean
                     Mean
                                       Mean
                                                        Mean
##
    3rd Qu.: 92.0
                     3rd Qu.:16.50
                                       3rd Qu.:31.00
                                                        3rd Qu.:10830
##
    Max.
            :100.0
                     Max.
                             :39.80
                                              :64.00
                                      Max.
                                                        Max.
                                                                :56233
##
      Grad.Rate
           : 10.00
   Min.
    1st Qu.: 53.00
##
```

```
## Median: 65.00
## Mean : 65.46
## 3rd Qu.: 78.00
## Max. :118.00

## Split
test_indices <- sample(1:nrow(college), 100)
college_test <- college[test_indices, ]
college_train <- college[-test_indices, ]

# Report the mean
mean_test <- mean(college_test$Outstate)
mean_train <- mean(college_train$Outstate)

mean_test

## [1] 9955.46</pre>
```

## [1] 10512.34

mean\_train

(a) Now, split your training data into two parts: validation data (100 observations), and the rest as estimation data. Use the random seed 4323.

```
# Split training data
validation_indices <- sample(1:nrow(college_train), 100)
college_validation <- college_train[validation_indices, ]
college_estimation <- college_train[-validation_indices, ]</pre>
```

(b) We're interested in predicting Enroll. First, run the linear regression on the estimation data including all variables. What is the feature variable with the highest p-value?

```
# Fit linear model
fit_all <- lm(Enroll ~ ., data = college_estimation)
summary(fit_all)</pre>
```

```
## PrivateYes -41.624379 32.844253 -1.267 0.205566
## Apps
              ## Accept
               0.195585
                         0.014194 13.779 < 2e-16 ***
## Top10perc
                                  3.711 0.000227 ***
               5.049148
                         1.360485
## Top25perc
              -2.906129
                         1.056250 -2.751 0.006127 **
## F.Undergrad
             ## P.Undergrad -0.018807
                         0.008618 -2.182 0.029507 *
## Outstate
              -0.005375
                         0.004428 -1.214 0.225312
## Room.Board
              -0.016068
                         0.011220 -1.432 0.152666
## Books
              -0.009285
                         0.056086 -0.166 0.868577
## Personal
               0.010304
                         0.014305
                                  0.720 0.471649
## PhD
              -0.269688
                         1.087562 -0.248 0.804245
## Terminal
              -0.360318
                        1.199981 -0.300 0.764083
## S.F.Ratio
               1.281315
                         3.066507
                                  0.418 0.676223
## perc.alumni
                         0.925684
                                   1.609 0.108100
               1.489767
## Expend
               0.006431
                         0.003387
                                   1.899 0.058080 .
## Grad.Rate
               0.790790
                         0.673169
                                   1.175 0.240603
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
## Residual standard error: 206 on 559 degrees of freedom
## Multiple R-squared: 0.9531, Adjusted R-squared: 0.9517
## F-statistic: 668.5 on 17 and 559 DF, p-value: < 2.2e-16
highest_p_value <- max(coef(summary(fit_all))[ , "Pr(>|t|)"])
highest_p_variable <- names(which(coef(summary(fit_all))[ , "Pr(>|t|)"] == highest_p_value))
highest_p_variable
```

## [1] "Books"

(c) Run the regression again, but this time, without that variable (with the highest p-value from previous regression) and find the feature variable with the highest p-value with the highest p-value in the new regression. Repeat this step until all the variables have p-value less than 0.1.

```
#first we exclude the books and enrolls
current_variables <- names(college_estimation)
current_variables <- current_variables[current_variables != "Enroll"]

current_variables <- setdiff(current_variables, "Books")

str(college$Private)</pre>
```

## Factor w/ 2 levels "No", "Yes": 2 2 2 2 2 2 2 2 2 ...

```
# Function to fit the model then get the variable with the highest p
get_highest_p_value_variable <- function(data, variables) {
  formula <- as.formula(paste("Enroll ~", paste(variables, collapse = " + ")))
  fit <- lm(formula, data = data)
  p_values <- summary(fit)$coefficients[-1, "Pr(>|t|)"]

if (length(p_values) == 0) return(list(variable = NULL, max_p_value = NA, model = fit))
```

```
max_p_value <- max(p_values)</pre>
  highest_p_variable <- names(which(p_values == max_p_value))
  # Handle the factor 'Private' correctly, it gives error because of the previous factorizeing function
  if ("PrivateYes" %in% highest_p_variable) {
    highest_p_variable <- "Private"
  }
  return(list(variable = highest_p_variable, max_p_value = max_p_value, model = fit))
max_iterations <- length(current_variables)</pre>
iterations <- 0
repeat {
 result <- get_highest_p_value_variable(college_estimation, current_variables)
  if (is.null(result$variable) || result$max_p_value <= 0.1) {</pre>
    final_model <- result$model</pre>
    break
  } else {
    current_variables <- setdiff(current_variables, result$variable)</pre>
    print(paste("Removing variable:", result$variable)) # Debugging print to see progress
  iterations <- iterations + 1
  if (iterations > max_iterations) {
    print("Reached maximum iterations.")
    break
  }
}
## [1] "Removing variable: PhD"
## [1] "Removing variable: S.F.Ratio"
## [1] "Removing variable: Personal"
## [1] "Removing variable: Terminal"
## [1] "Removing variable: Grad.Rate"
## [1] "Removing variable: Private"
## [1] "Removing variable: Room.Board"
# Summary of the final model
summary(final_model)
##
## Call:
## lm(formula = formula, data = data)
## Residuals:
                  1Q
                      Median
        Min
                                     3Q
## -1341.15 -61.55
                       -13.05
                                  52.49 1632.88
##
## Coefficients:
##
                 Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) 109.968345 39.697446
                                  2.770 0.00579 **
## Apps
              -0.051656  0.007847  -6.583  1.05e-10 ***
## Accept
              ## Top10perc
              5.270817
                         1.340517
                                   3.932 9.47e-05 ***
## Top25perc
              -3.057038
                        1.027394 -2.976 0.00305 **
## F.Undergrad 0.130682 0.004449 29.371
                                         < 2e-16 ***
## P.Undergrad -0.021442 0.008365 -2.563 0.01063 *
## Outstate
              -0.009877
                         0.003537 -2.793 0.00540 **
## perc.alumni 1.677391
                         0.879825
                                   1.907 0.05709 .
## Expend
               0.005032
                         0.002964
                                   1.698 0.09006 .
## ---
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' ' 1
## Residual standard error: 205.7 on 567 degrees of freedom
## Multiple R-squared: 0.9526, Adjusted R-squared: 0.9519
## F-statistic: 1267 on 9 and 567 DF, p-value: < 2.2e-16
```

(d) Find validation MSE of all the models in (b) and (c). Report the model with the smallest validation MSE.

```
calculate_mse <- function(model, data) {
  predictions <- predict(model, newdata = data)
  mse <- mean((data$Enroll - predictions)^2)
  return(mse)
}

# model with all variables
mse_all <- calculate_mse(fit_all, college_validation)

# final model from sec c
mse_final <- calculate_mse(final_model, college_validation)</pre>
mse_all
```

## [1] 33315.36

```
mse_final
```

## [1] 33588.9

```
best_model <- ifelse(mse_final < mse_all, "Final Model from sec c", " Model with All Variables")
best_model</pre>
```

## [1] " Model with All Variables"

(e) Report your test MSE of your chosen model in part (d).

```
chosen_model <- if(best_model == "Final Model from sec c") final_model else fit_all

test_mse <- calculate_mse(chosen_model, college_test)
test_mse</pre>
```

## [1] 46148.68

#### Question 3 (k-NN)

This question should be answered using the Carseats data set form ISLR2 package.

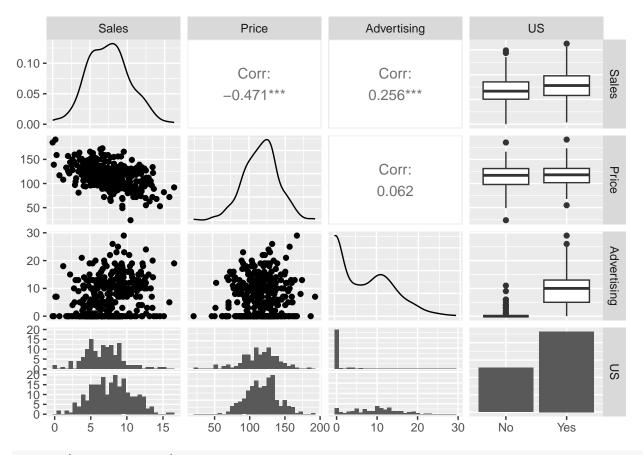
Make sure all categorical variables are set as factor variables, and omit any missing data.

```
data("Carseats")
str(Carseats)
##
   'data.frame':
                    400 obs. of 11 variables:
##
    $ Sales
                         9.5 11.22 10.06 7.4 4.15 ...
                 : num
                         138 111 113 117 141 124 115 136 132 132 ...
##
    $ CompPrice
                 : num
##
    $ Income
                         73 48 35 100 64 113 105 81 110 113 ...
                  : num
##
                         11 16 10 4 3 13 0 15 0 0 ...
    $ Advertising: num
##
    $ Population : num
                         276 260 269 466 340 501 45 425 108 131 ...
##
   $ Price
                         120 83 80 97 128 72 108 120 124 124 ...
                  : num
##
    $ ShelveLoc : Factor w/ 3 levels "Bad", "Good", "Medium": 1 2 3 3 1 1 3 2 3 3 ...
##
                  : num 42 65 59 55 38 78 71 67 76 76 ...
   $ Age
   $ Education : num 17 10 12 14 13 16 15 10 10 17 ...
##
                  : Factor w/ 2 levels "No", "Yes": 2 2 2 2 2 1 2 2 1 1 ...
    $ Urban
                  : Factor w/ 2 levels "No", "Yes": 2 2 2 2 1 2 1 2 1 2 ...
##
    $ US
# Convert categorical variables
Carseats$ShelveLoc <- as.factor(Carseats$ShelveLoc)</pre>
Carseats$Urban <- as.factor(Carseats$Urban)</pre>
Carseats$US <- as.factor(Carseats$US)</pre>
summary(Carseats)
##
        Sales
                        CompPrice
                                         Income
                                                       Advertising
##
    Min.
           : 0.000
                      Min.
                             : 77
                                    Min.
                                            : 21.00
                                                              : 0.000
##
    1st Qu.: 5.390
                      1st Qu.:115
                                    1st Qu.: 42.75
                                                      1st Qu.: 0.000
##
   Median : 7.490
                      Median:125
                                    Median: 69.00
                                                      Median : 5.000
           : 7.496
##
    Mean
                      Mean
                             :125
                                    Mean
                                            : 68.66
                                                      Mean
                                                              : 6.635
##
    3rd Qu.: 9.320
                      3rd Qu.:135
                                    3rd Qu.: 91.00
                                                      3rd Qu.:12.000
##
```

```
:175
                                             :120.00
                                                               :29.000
    Max.
           :16.270
                      Max.
                                     Max.
                                                       Max.
##
      Population
                         Price
                                       ShelveLoc
                                                          Age
                                                                        Education
##
                                             : 96
                                                            :25.00
                                                                             :10.0
    Min.
           : 10.0
                             : 24.0
                                      Bad
                                                    Min.
                                                                     Min.
                     Min.
    1st Qu.:139.0
                     1st Qu.:100.0
                                      Good : 85
                                                    1st Qu.:39.75
##
                                                                     1st Qu.:12.0
##
    Median :272.0
                     Median :117.0
                                      Medium:219
                                                    Median :54.50
                                                                     Median:14.0
##
    Mean
            :264.8
                     Mean
                             :115.8
                                                    Mean
                                                            :53.32
                                                                     Mean
                                                                             :13.9
##
    3rd Qu.:398.5
                                                    3rd Qu.:66.00
                                                                      3rd Qu.:16.0
                     3rd Qu.:131.0
##
    Max.
            :509.0
                     Max.
                             :191.0
                                                    Max.
                                                            :80.00
                                                                     Max.
                                                                             :18.0
##
    Urban
                 US
##
    No :118
              No :142
##
    Yes:282
              Yes:258
##
##
##
##
```

(a) Set 10% of whole data as a test set, and the rest as a training set. Split the training set into validation set (10% of training data) and the rest of the training set as a estimation set. Use the random seed 4324.

```
set.seed(4324)
# Split data
n <- nrow(Carseats)</pre>
test_indices <- sample(1:n, size = round(0.1 * n))</pre>
carseats_test <- Carseats[test_indices, ]</pre>
carseats_train <- Carseats[-test_indices, ]</pre>
train_indices <- sample(1:nrow(carseats_train), size = round(0.9 * nrow(carseats_train)))</pre>
carseats_estimation <- carseats_train[train_indices, ]</pre>
carseats_validation <- carseats_train[-train_indices, ]</pre>
# Check the sizee
cat("Test set size:", nrow(carseats_test), "\n")
## Test set size: 40
cat("Training set size:", nrow(carseats_train), "\n")
## Training set size: 360
cat("Estimation set size:", nrow(carseats_estimation), "\n")
## Estimation set size: 324
cat("Validation set size:", nrow(carseats_validation), "\n")
## Validation set size: 36
 (b) Conduct the EDA on the training set.
ggpairs(carseats_train, columns = c("Sales", "Price", "Advertising", "US"))
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



summary(carseats\_train)

```
##
        Sales
                         CompPrice
                                             Income
                                                            Advertising
##
    {\tt Min.}
            : 0.000
                       Min.
                              : 77.0
                                        {\tt Min.}
                                                : 21.00
                                                           Min.
                                                                   : 0.000
    1st Qu.: 5.345
                                                           1st Qu.: 0.000
##
                       1st Qu.:115.0
                                        1st Qu.: 43.75
    Median : 7.510
                       Median :125.0
                                        Median : 69.00
                                                           Median : 5.000
##
##
    Mean
            : 7.495
                       Mean
                              :125.3
                                        Mean
                                                : 68.95
                                                           Mean
                                                                   : 6.575
    3rd Qu.: 9.320
                                        3rd Qu.: 91.00
##
                       3rd Qu.:135.0
                                                           3rd Qu.:11.250
##
    Max.
            :16.270
                       Max.
                               :175.0
                                        Max.
                                                :120.00
                                                           Max.
                                                                   :29.000
##
      Population
                          Price
                                        ShelveLoc
                                                           Age
                                                                         Education
##
    Min.
            : 10.0
                             : 24.0
                                       Bad
                                              : 85
                                                             :25.00
                                                                               :10.00
                     Min.
                                                     Min.
                                                                       Min.
    1st Qu.:136.2
                     1st Qu.:100.8
                                       Good : 77
                                                      1st Qu.:39.00
                                                                       1st Qu.:12.00
##
##
    Median :268.5
                     Median :118.0
                                       Medium:198
                                                     Median :54.00
                                                                       Median :14.00
##
            :262.7
                              :116.4
                                                             :52.79
                                                                       Mean
                                                                               :13.91
    Mean
                     Mean
                                                     Mean
    3rd Qu.:393.2
                                                      3rd Qu.:65.00
                                                                       3rd Qu.:16.00
##
                     3rd Qu.:131.0
##
    Max.
            :509.0
                              :191.0
                                                             :80.00
                                                                               :18.00
                     Max.
                                                     Max.
                                                                       Max.
##
    Urban
                 US
               No :128
##
    No :109
    Yes:251
               Yes:232
##
##
##
##
##
```

(c) We're going to fit linear regression models to predict Sales using Price, US, and Advertising.

```
# Define the models
models <- list(</pre>
  model1 = Sales ~ Price,
  model2 = Sales ~ US,
  model3 = Sales ~ Advertising,
  model4 = Sales ~ Price + US,
  model5 = Sales ~ US + Advertising,
 model6 = Sales ~ Price + Advertising,
 model7 = Sales ~ Price + US + Advertising,
 model8 = Sales ~ Price * US + Advertising,
  model9 = Sales ~ Price * Advertising + US * Advertising
)
# Fit the models
fitted_models <- lapply(models, function(formula) {</pre>
  lm(formula, data = carseats_estimation)
})
lapply(fitted_models, summary)
## $model1
##
## Call:
## lm(formula = formula, data = carseats_estimation)
## Residuals:
##
       Min
                1Q Median
## -6.5020 -1.7350 -0.1111 1.5547 7.4694
## Coefficients:
                Estimate Std. Error t value Pr(>|t|)
## (Intercept) 13.910602
                           0.685815 20.283
                                              <2e-16 ***
## Price
               -0.055544
                           0.005768 -9.629
                                              <2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
## Residual standard error: 2.509 on 322 degrees of freedom
## Multiple R-squared: 0.2236, Adjusted R-squared: 0.2212
## F-statistic: 92.72 on 1 and 322 DF, p-value: < 2.2e-16
##
##
## $model2
##
## Call:
## lm(formula = formula, data = carseats_estimation)
## Residuals:
##
       Min
                1Q Median
## -7.4229 -1.9254 -0.0779 1.7318 8.4771
## Coefficients:
##
               Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)
                6.8199
                           0.2608
                                    26.15 < 2e-16 ***
## USYes
                0.9730
                           0.3255
                                     2.99 0.00301 **
## ---
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' 1
## Residual standard error: 2.809 on 322 degrees of freedom
## Multiple R-squared: 0.02701,
                                 Adjusted R-squared: 0.02399
## F-statistic: 8.939 on 1 and 322 DF, p-value: 0.003007
##
##
## $model3
##
## Call:
## lm(formula = formula, data = carseats_estimation)
## Residuals:
##
      Min
               1Q Median
                               3Q
                                      Max
## -7.2011 -1.8981 -0.1626 1.6948 8.3439
##
## Coefficients:
##
              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 6.77262
                          0.21409
                                   31.63 < 2e-16 ***
## Advertising 0.10269
                                     4.50 9.5e-06 ***
                          0.02282
## Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
## Residual standard error: 2.762 on 322 degrees of freedom
## Multiple R-squared: 0.05917,
                                  Adjusted R-squared: 0.05625
## F-statistic: 20.25 on 1 and 322 DF, p-value: 9.495e-06
##
##
## $model4
##
## Call:
## lm(formula = formula, data = carseats_estimation)
## Residuals:
##
               1Q Median
      Min
                               3Q
                                      Max
## -6.9459 -1.6341 -0.0128 1.5022 6.9825
##
## Coefficients:
               Estimate Std. Error t value Pr(>|t|)
##
                         0.679886 19.661 < 2e-16 ***
## (Intercept) 13.366986
## Price
              -0.057600
                          0.005639 -10.215 < 2e-16 ***
## USYes
               1.219740
                          0.284194
                                    4.292 2.35e-05 ***
## ---
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' 1
## Residual standard error: 2.444 on 321 degrees of freedom
## Multiple R-squared: 0.2657, Adjusted R-squared: 0.2611
## F-statistic: 58.08 on 2 and 321 DF, p-value: < 2.2e-16
##
##
## $mode15
```

```
##
## Call:
## lm(formula = formula, data = carseats estimation)
## Residuals:
##
     Min
             1Q Median
                           3Q
                                 Max
## -7.201 -1.898 -0.163 1.695 8.344
##
## Coefficients:
##
               Estimate Std. Error t value Pr(>|t|)
## (Intercept) 6.772970
                          0.257207 26.333 < 2e-16 ***
                                   -0.002 0.99804
              -0.001071
                          0.434975
## Advertising 0.102746
                          0.031016
                                    3.313 0.00103 **
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.766 on 321 degrees of freedom
## Multiple R-squared: 0.05917,
                                   Adjusted R-squared: 0.05331
## F-statistic: 10.09 on 2 and 321 DF, p-value: 5.602e-05
##
## $model6
##
## Call:
## lm(formula = formula, data = carseats_estimation)
## Residuals:
      Min
               1Q Median
                               3Q
                                      Max
## -7.9029 -1.6211 -0.0784 1.4224 6.2348
##
## Coefficients:
##
               Estimate Std. Error t value Pr(>|t|)
## (Intercept) 13.460354
                          0.653946 20.583 < 2e-16 ***
              -0.058477
                          0.005486 -10.659 < 2e-16 ***
## Price
## Advertising 0.121005
                          0.019717
                                    6.137 2.48e-09 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.377 on 321 degrees of freedom
## Multiple R-squared: 0.3051, Adjusted R-squared: 0.3008
## F-statistic: 70.47 on 2 and 321 DF, p-value: < 2.2e-16
##
##
## $model7
## Call:
## lm(formula = formula, data = carseats_estimation)
## Residuals:
##
      Min
               1Q Median
                               ЗQ
## -7.8765 -1.5884 -0.0729 1.4541 6.2784
## Coefficients:
##
               Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) 13.422801
                         0.662411 20.264 < 2e-16 ***
## Price
                         0.005497 -10.651 < 2e-16 ***
             -0.058551
## USYes
                         0.374562
              0.140688
                                  0.376
                                            0.707
## Advertising 0.114246
                         0.026713
                                  4.277 2.51e-05 ***
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' 1
## Residual standard error: 2.38 on 320 degrees of freedom
## Multiple R-squared: 0.3054, Adjusted R-squared: 0.2989
## F-statistic: 46.9 on 3 and 320 DF, p-value: < 2.2e-16
##
##
## $model8
##
## Call:
## lm(formula = formula, data = carseats_estimation)
##
## Residuals:
      Min
              1Q Median
                             30
                                    Max
## -7.8793 -1.5892 -0.0706 1.4578 6.2890
##
## Coefficients:
               Estimate Std. Error t value Pr(>|t|)
##
## (Intercept) 13.3845172 0.9776936 13.690 < 2e-16 ***
## Price
             0.2087377 1.3304978
                                   0.157
                                            0.875
## Advertising 0.1142863 0.0267656
                                    4.270 2.58e-05 ***
## Price:USYes -0.0005924 0.0111127 -0.053
                                            0.958
## Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' 1
## Residual standard error: 2.384 on 319 degrees of freedom
## Multiple R-squared: 0.3054, Adjusted R-squared: 0.2967
## F-statistic: 35.07 on 4 and 319 DF, p-value: < 2.2e-16
##
##
## $model9
##
## lm(formula = formula, data = carseats_estimation)
## Residuals:
      Min
              1Q Median
                             30
## -7.8611 -1.6028 -0.0162 1.4753 6.1163
## Coefficients:
##
                     Estimate Std. Error t value Pr(>|t|)
                   ## (Intercept)
## Price
                   -0.0611064 0.0073431
                                        -8.322 2.59e-15 ***
## Advertising
                   -0.1149719
                              0.1694489
                                         -0.679
                                                  0.498
## USYes
                    0.0159397 0.3939262
                                         0.040
                                                  0.968
## Price: Advertising 0.0004444 0.0008204
                                         0.542
                                                  0.588
## Advertising:USYes 0.1819881 0.1464572
                                         1.243
                                                  0.215
## ---
```

```
## Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.381 on 318 degrees of freedom
## Multiple R-squared: 0.3096, Adjusted R-squared: 0.2987
## F-statistic: 28.52 on 5 and 318 DF, p-value: < 2.2e-16
Candidate models:
model 1: Sales~Price
model 2: Sales~US
model 3: Sales~Advertising
model 4: Sales~Price+US
model 5: Sales~US+Advertising
model 6: Sales~Price+Advertising
model 7: Sales~Price+US+Price*US
model 8: Sales~US+Advertising+US*Advertising
model 9: Sales~Price+Advertising+Price*Advertising
Store all regression models in one list. Run the regressions on the estimation data.
```

(e) Calculate validation MSE of all models. Choose a single model with the lowest validation MSE.

```
predictions <- predict(model, newdata = data)</pre>
 mse <- mean((data$Sales - predictions)^2)</pre>
  return(mse)
validation_mse <- sapply(fitted_models, calculate_mse, data = carseats_validation)</pre>
validation_mse
     model1
              model2
                        model3
                                  model4
                                           model5
                                                     model6
                                                               model7
                                                                        model8
## 6.794518 7.796904 7.556583 6.190823 7.556891 6.072983 6.050427 6.049539
     model9
## 6.305631
# Identify lowest validation MSE
best_model_index <- which.min(validation_mse)</pre>
best_model <- fitted_models[[best_model_index]]</pre>
cat("Best model is:", names(models)[best_model_index], "with MSE =", validation_mse[best_model_index],
## Best model is: model8 with MSE = 6.049539
```


(f) Report your test MSE. Provide a scatter plot of predicted Sales and observed Sales of the test data.

```
test_mse <- calculate_mse(best_model, carseats_test)</pre>
cat("Test MSE for the best model:", test_mse, "\n")
```

## Test MSE for the best model: 6.452404

calculate\_mse <- function(model, data) {</pre>

### **Predicted vs Observed Sales**



#### Question 4 (k-NN and decision tree)

This question relates to the Boston data set of ISLR2 package.

```
library(caret)

## Warning: package 'caret' was built under R version 4.3.3

## Loading required package: lattice

library(rpart)
library(rpart.plot)

## Warning: package 'rpart.plot' was built under R version 4.3.3

set.seed(432)

trn.idx=sample(1:nrow(ISLR2::Boston),450)

tst.boston=ISLR2::Boston[-trn.idx,]

trn.boston=ISLR2::Boston[trn.idx,]
We are splitting the data into two parts: a testing data that contains 56 observations, and the rest 450 observations as training data.
```

- The goal is to model medv (our response variable) with all the other variables in the data.
- In this HW, we'll not worry about scaling variables. We'll tackle that in the future.
- (a) Use the following validation-estimation split.

```
set.seed(1)
val.idx=sample(1:nrow(trn.boston),45)
val.boston=trn.boston[val.idx,]
est.boston=trn.boston[-val.idx,]
# Check the sizes of the datasets
cat("Training set size:", nrow(trn.boston), "\n")

## Training set size: 450

cat("Estimation set size:", nrow(est.boston), "\n")

## Estimation set size: 405

cat("Validation set size:", nrow(val.boston), "\n")

## Validation set size: 45


cat("Test set size:", nrow(tst.boston), "\n")

## Test set size: 56
```

- Use the estimation data and knnreg function of caret package to perform KNN.
- Train KNN models using values of k from 1 to 100 and calculate validation MSE for each k.
- Plot the validation MSE versus k and show them in the same graph.

```
train_x <- est.boston[, -which(names(est.boston) == "medv")]</pre>
train_y <- est.boston$medv</pre>
val_x <- val.boston[, -which(names(val.boston) == "medv")]</pre>
val_y <- val.boston$medv</pre>
# Initialize
k_values <- 1:100
validation_mse <- numeric(length(k_values))</pre>
for (k in k_values) {
  knn_model <- knnreg(train_x, train_y, k = k)</pre>
  val_predictions <- predict(knn_model, val_x)</pre>
  validation_mse[k] <- mean((val_y - val_predictions)^2)</pre>
}
# Plot
plot(k_values, validation_mse, type = "b", col = "blue", pch = 19,
     xlab = "Number of Neighbors (k)", ylab = "Validation MSE",
     main = "Validation MSE vs. Number of Neighbors (k)")
```

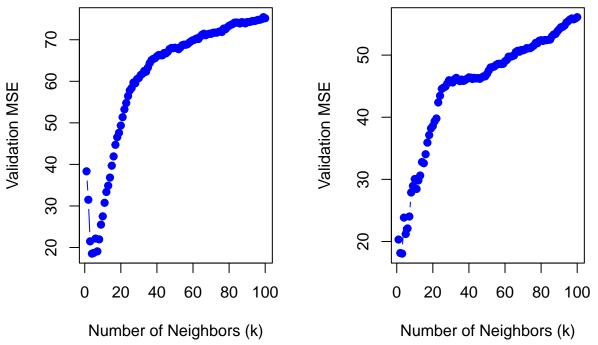
### **Validation MSE vs. Number of Neighbors (k)**



(b) Repeat (a) with different random seeds, (2,3), and see if your answer changes. If so, why does it change?

```
knn_analysis_with_seed <- function(seed) {</pre>
  set.seed(seed)
  val.idx <- sample(1:nrow(trn.boston), 45)</pre>
  val.boston <- trn.boston[val.idx, ]</pre>
  est.boston <- trn.boston[-val.idx, ]</pre>
  train_x <- est.boston[, -which(names(est.boston) == "medv")]</pre>
  train y <- est.boston$medv</pre>
  val_x <- val.boston[, -which(names(val.boston) == "medv")]</pre>
  val_y <- val.boston$medv</pre>
  validation_mse <- numeric(length(k_values))</pre>
  for (k in k_values) {
    knn_model <- knnreg(train_x, train_y, k = k)</pre>
    val_predictions <- predict(knn_model, val_x)</pre>
    validation_mse[k] <- mean((val_y - val_predictions)^2)</pre>
  return(validation_mse)
#for other seeds
seeds \leftarrow c(2, 3)
mse_results <- lapply(seeds, knn_analysis_with_seed)</pre>
# Plot the results
par(mfrow = c(1, 2))
for (i in 1:length(seeds)) {
  plot(k_values, mse_results[[i]], type = "b", col = "blue", pch = 19,
       xlab = "Number of Neighbors (k)", ylab = "Validation MSE",
       main = paste("Validation MSE vs. k (Seed =", seeds[i], ")"))
}
```

# Validation MSE vs. k (Seed = 2) Validation MSE vs. k (Seed = 3)



Answer: The validation MSE changes with different random seeds becuz each seed results in a diffrent random split of the data. This means that the estimation and validation sets are different, which effects how the kNN model is trained and evaluted. The differences in data subsets lead to variations in model performance and validation MSE. This variablity is typical in machine learning and show the importance of using consistent data splits or multiple runs to assess model stablity.


- (c) Use the estimation/validation data from (a) with random seed (1) and rpart and rpart.plotfunction to perform decision tree.
  - Start with default setting of R.
- Train decision tree models using cp=0, 0.001, 0.01, 0.1.
- Students may explore other tuning parameters as needed.
- Show your tree results using rpart.plot function.
- Compute validation MSE versus different cp values.
- Choose cp with lowest validation MSE.

```
cp_values <- c(0, 0.001, 0.01, 0.1)

models <- list()
validation_mse <- numeric(length(cp_values))

# Train models with diff cp
for (i in 1:length(cp_values)) {
    cp <- cp_values[i]
    model <- rpart(medv ~ ., data = est.boston, control = rpart.control(cp = cp))
    models[[i]] <- model
    val_predictions <- predict(model, newdata = val.boston)
    validation_mse[i] <- mean((val.boston$medv - val_predictions)^2)</pre>
```

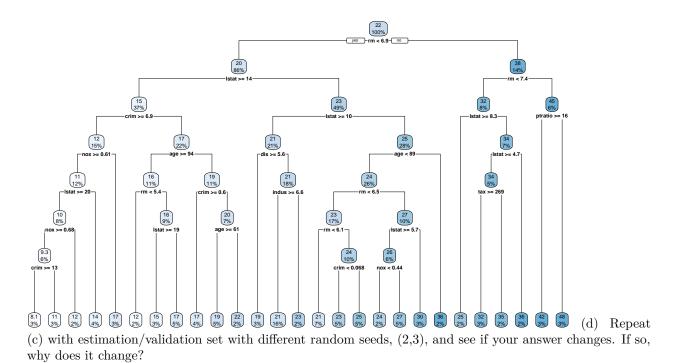
## Validation MSE vs. cp



```
# Identify and print the best cp value
best_model_index <- which.min(validation_mse)
best_cp <- cp_values[best_model_index]
cat("The best cp value is", best_cp, "with the lowest validation MSE of", validation_mse[best_model_ind</pre>
```

## The best cp value is 0.001 with the lowest validation MSE of 20.72719

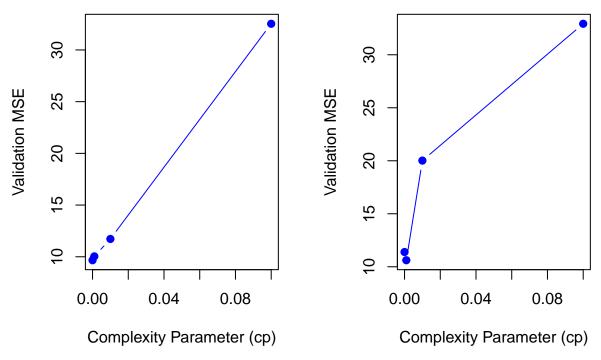
```
# Display the tree
best_model <- models[[best_model_index]]
rpart.plot(best_model)</pre>
```



```
decision_tree_analysis_with_seed <- function(seed) {</pre>
  set.seed(seed)
  val.idx <- sample(1:nrow(trn.boston), 45)</pre>
  val.boston <- trn.boston[val.idx, ]</pre>
  est.boston <- trn.boston[-val.idx, ]</pre>
  validation_mse <- numeric(length(cp_values))</pre>
  for (i in 1:length(cp_values)) {
    cp <- cp_values[i]</pre>
    model <- rpart(medv ~ ., data = est.boston, control = rpart.control(cp = cp))</pre>
    val_predictions <- predict(model, newdata = val.boston)</pre>
    validation_mse[i] <- mean((val.boston$medv - val_predictions)^2)</pre>
    # Print the MSE for each cp
    cat("Seed =", seed, "cp =", cp, "- Validation MSE:", validation_mse[i], "\n")
  # Identify the best cp
  best_cp_index <- which.min(validation_mse)</pre>
  best_cp <- cp_values[best_cp_index]</pre>
  best_mse <- validation_mse[best_cp_index]</pre>
  cat("For seed =", seed, "the best cp value is", best_cp, "with the lowest validation MSE of", best_ms
  return(validation_mse)
}
# Calculate val MSEs for diff seeds
seeds \leftarrow c(2, 3)
mse_results <- lapply(seeds, decision_tree_analysis_with_seed)</pre>
```

```
## Seed = 2 \text{ cp} = 0 - \text{Validation MSE}: 9.666538
## Seed = 2 cp = 0.001 - Validation MSE: 10.03971
## Seed = 2 cp = 0.01 - Validation MSE: 11.72382
## Seed = 2 \text{ cp} = 0.1 - \text{Validation MSE}: 32.53138
## For seed = 2 the best cp value is 0 with the lowest validation MSE of 9.666538
## Seed = 3 cp = 0 - Validation MSE: 11.39089
## Seed = 3 cp = 0.001 - Validation MSE: 10.61798
## Seed = 3 cp = 0.01 - Validation MSE: 20.01785
## Seed = 3 \text{ cp} = 0.1 - \text{Validation MSE}: 32.92429
## For seed = 3 the best cp value is 0.001 with the lowest validation MSE of 10.61798
# Plot the results
par(mfrow = c(1, 2))
for (i in 1:length(seeds)) {
  plot(cp_values, mse_results[[i]], type = "b", col = "blue", pch = 19,
       xlab = "Complexity Parameter (cp)", ylab = "Validation MSE",
       main = paste("Validation MSE vs. cp (Seed =", seeds[i], ")"))
}
```

## Validation MSE vs. cp (Seed = 2 Validation MSE vs. cp (Seed = 3



In this ML methods such as the KNN the results changed because of having different testing data. As you can see even the best cp can be different for different random seed numbers as in 1 and 3 the best is cp=0.001 but in the random see =2 it has ben the cp=0. However, you can see the change between the cp=0.001 and cp=0 is not a lot so it can change by changing the random seed.